

Course Overview

CS230: System Programming
1st Lecture, Aug. 31, 2021

Instructors:

Jongse Park

Jongse Park



■ Background

- Assistant professor, CASys lab, KAIST (2019 – present)
- Deep learning acceleration system architect, Bigstream (2018 – 2019)
- Ph.D. from ACT Lab, Georgia Institute of Technology (2018)
- M.S. from CASys lab, KAIST (2012)
- B.S. from Sogang University (2010)

■ Contact

- E3-1, Room 4403
- jspark@casys.kaist.ac.kr
- <https://jongse-park.github.io/>

■ Research Interests

- Computer Architecture
- Systems for Artificial Intelligence
- Edge-to-cloud computing

Overview

- Course theme
- Academic integrity
- Course logistics
- Course schedule
- Five realities

COURSE THEME

Computing in every aspect of our lives



Home Automation



Bioinformatics



Sharing Economy



Ecommerce



Online Education



Precision Agriculture



Computational Finance

System: Foundation of All Innovations


All innovations in Computer Science




Computer Systems

Example: Huge Success of Deep Learning

Applications




Algorithms



System



V100



DGX-2

This course is

- CS230 **System** Programming
- The course is an “introduction to computer systems” from the perspective of programmers
- If you expected System **Programming**, this is not a right course.

Course Perspective

- Most Systems Courses are Builder-Centric
 - Computer Architecture
 - Design pipelined processor
 - Operating Systems
 - Implement sample portions of operating system
 - Compilers
 - Write compiler for simple language
 - Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is Programmer-Centric

- Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
 - **We bring out the hidden hacker in everyone!**

ACADEMIC INTEGRITY

Cheating: Description

■ Pay Close Attention

■ What is cheating?

- Sharing code: by copying, retyping, **looking at**, or supplying a file
- Describing: verbal description of code from one person to another.
- Coaching: helping your friend to write a lab, line by line
- **Searching the Web for solutions**
- Copying code from a previous course or online solution
 - You are only allowed to use code we supply, or from the CS:APP website

■ What is NOT cheating?

- Explaining how to use systems or tools
- Helping others with high-level design issues

■ Ignorance is not an excuse

Cheating: Consequences

- Penalty for cheating:
 - Minimum penalty: one letter grade downgrade (A0 → B0)
 - Possible removal from course with failing grade (F)
 - Your instructors' personal contempt
- Detection of cheating:
 - We have sophisticated tools for detecting code plagiarism
 - In the prior semester, >30 students were caught cheating
- Don't do it!
 - Start early
 - Ask the staff for help when you get stuck

COURSE LOGISTICS

Live video lecture via Zoom

- Join Zoom Meeting
 - <https://kaist.zoom.us/j/88319898349?pwd=NVVya2RSMmNERWxxbnBISVQvYIV2UT09>
- Meeting ID: 883 1989 8349
- Passcode: 083121


Textbooks

■ Randal E. Bryant and David R. O'Hallaron,

- *Computer Systems: A Programmer's Perspective, Third Edition* (CS:APP3e), Pearson, 2016
- <http://csapp.cs.cmu.edu>
- Highly recommend the original English version (you'll have to get used to textbooks in English as soon as possible to survive in CS)
- This book really matters for the course! (for labs and exams)
- **Reading the textbook is NOT optional, but required**
 - I will post required sections in KLMS.
 - There will be occasional quizzes about the reading assignments.

■ Optional textbook: Brian Kernighan and Dennis Ritchie,

- *The C Programming Language*, Second Edition, Prentice Hall, 1988
- Still the best book about C, from the originators



Course Components

- Lectures
 - Higher level concepts
- Labs (7 assignments)
 - **The heart of the course**
 - About 2 weeks each
 - Provide in-depth understanding of an aspect of systems
 - Programming and measurement
- Exams (midterm + final)
 - Test your understanding of concepts & mathematical principles

Getting Help

- Class web page: KLMS
 - Complete schedule of lectures, exams, and assignments
 - Copies of lectures, assignments, exams, solutions
 - **Update your email address (Important announcements can be made via emails and KLMS postings)**
- Q&A board: Piazza
 - Clarifications to assignments
 - Signup link: piazza.com/kaist.ac.kr/fall2021/cs230
 - Access code: cs230casys
 - **Students are also encouraged to answer questions**

Getting Help

- Staff mailing list: **cs230b_ta@casys.kaist.ac.kr**
 - Use this for all communication with the teaching staff
 - Always CC staff mailing list during email exchanges
- Office hours or TA sessions by TAs
 - The schedule will be announced on KLMS

Policies: Labs And Exams

■ Work groups

- You must **work alone** on all lab assignments

■ Handins

- Labs due at 11:59pm on Tues or Thurs
- Electronic handins (no exceptions!)

■ Exams

- Midterm + Final (traditional exams)
- **At least 2/3 of exam questions are related to the lab assignments**

■ Appealing grades

- In **writing** to Prof. Jongse Park within 7 days of completion of grading

Timeliness

- Lateness penalties
 - Get penalized **30%**
 - No handins later than **1 day after due date**
- Catastrophic events
 - Major illness, death in family, ...
 - Formulate a plan (with your academic advisor) to get back on track
- Advice
 - Once you start running late, it's really hard to catch up

Policies: Grading

- Exams (54%):
 - midterm (27%), final (27%)
- Labs (40%): weighted according to effort
 - Labs + quiz
 - There will be occasional quizzes
- Attendance (6%)
 - We will use KLMS for checking
 - Attendance of the first week (8/31 and 9/1) will not be checked as the enrollment changes during the period
 - If you miss 1/3 of class meetings (9 class sessions), the final grade will be automatically F.

COURSE SCHEDULE

Labs

- *Tentative* lab schedule
- Lab 1: Linked list (Tuesday, 9/14)
- Lab 2: Data lab (Thursday, 9/30)
- Lab 3: Bomb lab (Tuesday, 10/12)
- Lab 4: Attack lab (Thursday, 10/28)
- Lab 5: Tsh lab (Thursday, 11/11)
- Lab 6: Malloc lab (Thursday, 11/25)
- Lab 7: Proxy lab (Thursday, 12/7)

Lab 1: Linked List

- The first lab for getting used to C and programming environments
- Due date: 9/14
- Implement simple functions for manipulating doubly-linked list data
 - Insert, delete, search
- Need to understand the basics of C, including pointers

Programs and Data

■ Topics

- Bits operations, arithmetic, assembly language programs
- Representation of C control and data structures
- Includes aspects of architecture and compilers

■ Assignments

- L2 (datalab): Manipulating bits
- L3 (bomblab): Defusing a binary bomb
- L4 (attacklab): The basics of code injection attacks

Exceptional Control Flow

■ Topics

- Hardware exceptions, processes, process control, Unix signals, nonlocal jumps
- Includes aspects of compilers, OS, and architecture

■ Assignments

- L5 (tshlab): Writing your own Unix shell.
 - A first introduction to concurrency

Virtual Memory

■ Topics

- Virtual memory, address translation, dynamic storage allocation
- Includes aspects of architecture and OS

■ Assignments

- L6 (malloclab): Writing your own malloc package
 - Get a real feel for systems-level programming

Networking, and Concurrency

■ Topics

- High level and low-level I/O, network programming
- Internet services, Web servers
- concurrency, concurrent server design, threads
- I/O multiplexing with select
- Includes aspects of networking, OS, and architecture

■ Assignments

- L7 (proxylab): Writing your own Web proxy
 - Learn network programming and more about concurrency and synchronization.

Concurrency

■ Topics

- Concurrency, concurrent server design, threads
- I/O multiplexing with select
- Includes aspects of networking, OS, and architecture

Lab Rationale

- Each lab has a well-defined goal such as solving a puzzle or winning a contest
- Doing the lab should result in new skills and concepts
- We try to use competition in a fun and healthy way
 - Set a reasonable threshold for full credit

FIVE REALITIES


Course Theme: Abstraction Is Good But Don't Forget Reality

- Most CS and CE courses emphasize abstraction
 - Abstract data types
 - Asymptotic analysis
- These abstractions have limits
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- Useful outcomes from taking CS230
 - Become more effective programmers
 - Able to and eliminate bugs efficiently
 - Able to understand and tune for program performance
 - Prepare for later “systems” classes in CS
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

Great Reality #1: Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

- Float's: Yes!



- Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

- Unsigned & Signed Int's: Yes!

- Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Computer Arithmetic

- **Does not generate random values**
 - Arithmetic operations have important mathematical properties
- **Cannot assume all “usual” mathematical properties**
 - Due to finiteness of representations
 - Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
 - Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs
- **Observation**
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ **Memory is not unbounded**

- It must be allocated and managed
- Many applications are memory dominated

■ **Memory referencing bugs especially pernicious**

- Effects are distant in both time and space

■ **Memory performance is not uniform**

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault


- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Great Reality #4: There's more to performance than asymptotic complexity

- Constant factors matter too!
- And even exact op count does not predict performance
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- Must understand system to optimize performance
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],  
           int dst[2048][2048])  
{  
    int i,j;  
    for (i = 0; i < 2048; i++)  
        for (j = 0; j < 2048; j++)  
            dst[i][j] = src[i][j];  
}
```

```
void copyji(int src[2048][2048],  
           int dst[2048][2048])  
{  
    int i,j;  
    for (j = 0; j < 2048; j++)  
        for (i = 0; i < 2048; i++)  
            dst[i][j] = src[i][j];  
}
```

4.3ms

2.0 GHz Intel Core i7 Haswell

81.8ms

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Great Reality #5: Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance
- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Welcome and Enjoy!