

Exceptional Control Flow: Exceptions and Processes

Entire system

CS230 System Programming
10th Lecture

Instructors:

Jongse Park

hit keyboard, outside interruption, ...

Today

- **Exceptional Control Flow**
- Exceptions
- Processes
- Process Control

Control Flow

- Processors do only one thing: *at a time*
- From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
- This sequence is the CPU's *control flow* (or *flow of control*)
trace of instructions
Physical control flow



Altering the Control Flow

■ Up to now: two mechanisms for changing control flow:

- Jumps and branches *within program*
- Call and return *→ more restricted control flow*
React to changes in *program state*
Condition codes, flags, ...

■ Insufficient for a useful system:

Difficult to react to changes in *system state*

- Data arrives from a disk or a network adapter *→ disk should notify*
- Instruction divides by zero *→ possible at runtime* *processor should react*
- User hits Ctrl-C at the keyboard
- System timer expires *→ at certain period of time*
Kill infinite loop *Create event → time is up!*

■ System needs mechanisms for “exceptional control flow”

Exceptional Control Flow

- Exists at all levels of a computer system

- Low level mechanisms

- 1. **Exceptions**

- Change in control flow in response to a system event
(i.e., change in system state)
 - Implemented using combination of hardware and OS software

- Higher level mechanisms

- 2. **Process context switch**

- Implemented by OS software and hardware timer

- 3. **Signals** *→ events created by software*

- Implemented by OS software

- 4. **Nonlocal jumps**: `setjmp()` and `longjmp()`

- Implemented by C runtime library




*→ jumping outside
of memory space*

Today

- Exceptional Control Flow
- Exceptions
- Processes
- Process Control


Exceptions



- An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)

Core

- Kernel is the memory-resident part of the OS
- Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



Exception Tables

Kernel Space

Exception numbers

Exception Table


0	•
1	•
2	•
...	
n-1	•

Code for exception handler 0

Code for exception handler 1

Code for exception handler 2

Code for exception handler n-1



- Each type of event has a unique exception number k
- $k = \text{index into exception table (a.k.a. interrupt vector)}$
- Handler k is called each time exception k occurs

Asynchronous Exceptions (Interrupts)

■ Caused by events external to the processor

- Indicated by setting the processor's *interrupt pin*
- Handler returns to "next" instruction



■ Examples:

- Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
- I/O interrupt from external device
 - Hitting Ctrl-C at the keyboard
 - Arrival of a packet from a network
 - Arrival of data from a disk

} not controlled by
processor



Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:

- Traps**



- Intentional
- Examples: **system calls**, breakpoint traps, special instructions
- Returns control to “next” instruction

communicate with kernel (printf)



- Faults**

- Unintentional but possibly recoverable
- Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
- Either re-executes faulting (“current”) instruction or aborts



- Aborts**

- Unintentional and unrecoverable
- Examples: illegal instruction, parity error, machine check
- Aborts current program never come back

System Calls

- Each x86-64 system call has a unique ID number
- Examples:

system ↗ for calling, you go to kernel space

Number	Name	Description
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

↳ number as identifier (for kernel)

System Call Example: Opening File


- User calls: `open (filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

```
00000000000e5d70 <__open>:
...
e5d79: b8 02 00 00 00    mov $0x2,%eax # open is syscall #2
e5d7e: 0f 05              syscall          # Return value in %rax
e5d80: 48 3d 01 f0 ff ff  cmp $0xffffffffffff001,%rax
...
e5dfa: c3                retq
...
```

as programmer, never initiate directly

syscall

get -> kernel



- `%rax` contains syscall number
 - Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
 - Return value in `%rax`
 - Negative value is an error corresponding to negative `errno`
- Failed*

Fault Example: Page Fault


- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];
main ()
{
    a[500] = 13;
}
```

properly allocated

80483b7:	c7 05 10 9d 04 08 0d	movl	\$0xd, 0x8049d10
----------	----------------------	------	------------------

*no guarantee @ [500] physical
location*




Fault Example: Invalid Memory Reference

```
int a[1000];
main ()
{
    a[5000] = 13;
}
```

hot in a+4x1000
memory

80483b7:	c7 05 60 e3 04 08 0d	movl \$0xd, 0x804e360
----------	----------------------	-----------------------



- ~~Segmentation fault~~ Sends SIGSEGV signal to user process
- User process exits with “segmentation fault”
not always recoverable
- ~~termination~~ kernel kills user program

Today

- Exceptional Control Flow
- Exceptions
- Processes
- Process Control

Processes

- Definition: A **process** is an instance of a running program.

- One of the most profound ideas in computer science
- Not the same as “program” or “processor”

virtual entity

↳ permanent physical entity

· /proj → multiple processes


- Process provides each program with two key abstractions:

- *Logical control flow* (conceptually)
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called *context switching*
- *Private address space*
 - Each program seems to have exclusive use of main memory.
 - Provided by kernel mechanism called *virtual memory*


Time sharing

Contiguously

physically support multi-processes



Multiprocessing: The Illusion



$P_i \rightarrow I'$ ve my own
CPU, Memory

- Computer runs many processes simultaneously (together)

- Applications for one or more users
 - Web browsers, email clients, editors, ...
- Background tasks
 - Monitoring network & I/O devices

time-sharing manner

Multiprocessing Example


```
xterm
Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads          11:47:07
Load Avg: 1.03, 1.13, 1.14 CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkededit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

PID   COMMAND    %CPU TIME    #TH  #WQ  #PORT #MREG RPRVT RSHRD RSIZE VPRVT VSIZE
99217- Microsoft Of 0.0 02:28.34 4    1    202   418   21M   24M   21M   66M   763M
99051  usbmuxd    0.0 00:04.10 3    1     47    66   436K  216K  480K  60M   2422M
99006  iTunesHelper 0.0 00:01.23 2    1     55    78   728K  3124K 1124K  43M   2429M
84286  bash        0.0 00:00.11 1    0     20    24   224K  732K  484K  17M   2378M
84285  xterm       0.0 00:00.83 1    0     32    73   656K  872K  692K  9728K 2382M
55939- Microsoft Ex 0.3 21:58.97 10   3    360   954   16M   65M   46M   114M  1057M
54751  sleep       0.0 00:00.00 1    0     17    20   92K   212K  360K  9632K 2370M
54739  launchdadd  0.0 00:00.00 2    1     33    50   488K  220K  1736K  48M   2409M
54737  top          6.5 00:02.53 1/1   0     30    29   1416K  216K  2124K  17M   2378M
54719  automountd  0.0 00:00.02 7    1     53    64   860K  216K  2184K  53M   2413M
54701  ocspd        0.0 00:00.05 4    1     61    54   1268K  2644K  3132K  50M   2426M
54661  Grab         0.6 00:02.75 6    3    222+  389+  15M+  26M+  40M+  75M+  2556M+
54659  cookied      0.0 00:00.15 2    1     40    61   3316K  224K  4088K  42M   2411M
53818  mdworker     0.0 00:01.67 4    1     52    91   7628K  7412K  16M   48M   2439M
50878  mdworker     0.0 00:11.17 3    1     57    91   2464K  6148K  9976K  44M   2434M
50410  mdworker     0.0 00:00.11 0    0     52    73   280K   872K  532K  9700K 2382M
50078  emacs        0.0 00:06.70 1    0     20    35   52K   216K  88K   18M   2392M
```

Running program “top” on Mac


- System has 123 processes, 5 of which are active
 - Identified by Process ID (PID)

Multiprocessing: The (Traditional) Reality




- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory

Multiprocessing: The (Traditional) Reality




- Save current registers in memory

Multiprocessing: The (Traditional) Reality



- Schedule next process for execution


Multiprocessing: The (Traditional) Reality



- Load saved registers and switch address space (context switch)



Multiprocessing: The (Modern) Reality




- **Multicore processors**
 - Multiple CPUs on single chip
 - Share main memory (and some of the caches)
 - Each can execute a separate process
 - Scheduling of processors onto cores done by kernel

Concurrent Processes

并发


- Each process is a logical control flow.
- Two processes *run concurrently* (are concurrent) if their flows overlap in time
- Otherwise, they are *sequential*
- Examples (running on single core):

- Concurrent: A & B, A & C
- Sequential: B & C



User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes as running in parallel with each other



Context Switching

Packed in DRAM
controls the processes

- Processes are managed by a shared chunk of memory-resident OS code called the *kernel* → core of Os
 - Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a *context switch*

processor's
top
PC

Process A

Process B

Time

(timer is up)

hardware creates interrupt

user code

kernel code

user code

kernel code

user code

} shares, saves registers
} context switch

} context switch



Today


- Exceptional Control Flow
- Exceptions
- Processes
- Process Control

System Call Error Handling

- ↙ fail here no problem, but HANDLE IT!*
- On error, Linux system-level functions typically return -1 and set global variable `errno` to indicate cause.
 - Hard and fast rule:
 - You must check the return status of every system-level function
 - Only exception is the handful of functions that return `void`
 - Example:

Krappener for error-check

```
if ((pid = fork()) < 0) {
    fprintf(stderr, "fork error: %s\n", strerror(errno));
    exit(0);
}
```



Error-reporting functions

when you call system-call

- Can simplify somewhat using an *error-reporting function*:

```
void unix_error(char *msg) /* Unix-style error */  
{  
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));  
    exit(0);  
}
```


```
if ((pid = fork()) < 0)  
    unix_error("fork error");
```

Error-handling Wrappers

- We simplify the code we present to you even further by using Stevens-style error-handling wrappers:

```
pid_t _Fork(void)
{
    pid_t pid;

    if ((pid = fork()) < 0)
        unix_error("Fork error");
    return pid;
}
```



pid = Fork();

~~ simple, protective

Obtaining Process IDs

- **pid_t getpid(void)**
 - Returns PID of current process

- **pid_t getppid(void)**
 - Returns PID of parent process

processes created by other processes
there's always parent process

Creating and Terminating Processes

From a programmer's perspective, we can think of a process as being in one of three states

■ Running

- Process is either executing, or waiting to be executed and will eventually be scheduled (i.e., chosen to execute) by the kernel

■ Stopped

- Process execution is *suspended* and will not be scheduled until further notice (next lecture when we study signals)

does not schedule
↑
 $\text{Ctrl} + Z$

■ Terminated

- Process is stopped permanently

kills the process

Terminating Processes

■ Process becomes terminated for one of three reasons:

- Receiving a signal whose default action is to terminate (next lecture)
- Returning from the `main` routine \rightsquigarrow I'm done
- Calling the `exit` function \rightsquigarrow system-call
sending kernel: gdp

■ `void exit(int status)`

- Terminates with an *exit status* of `status`
- Convention: normal return status is 0, nonzero on error
- Another way to explicitly set the *exit status* is to return an integer value from the `main` routine

■ `exit` is called **once** but **never** returns.

```
int main() {  
    }  
    return 0;
```


Creating Processes

- Parent process creates a new running child process by calling **fork**

OS creates shell for you (process for the first)

int fork(void)
function call

- Returns 0 to the child process, child's PID to parent process
- Child is almost identical to parent:
 - Child gets an identical (but separate) copy of the parent's virtual address space. (*same copy of data, stack*)
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent



- fork is interesting (and often confusing) because it is called **once** but returns **twice**

process { parent } → child's PID } return value
 child → 0

fork Example

```

int main()
{
    pid_t pid;
    int x = 1;
    2 processes running
    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }
    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}


```

fork.c

```

linux> ./fork
parent: x=0
child : x=2

```



- Call once, return twice
Context switching
- Concurrent execution
 - Can't predict execution order of parent and child
- Duplicate but separate address space
 - x has a value of 1 when fork returns in parent and child *Separate copy!*
 - Subsequent changes to x are independent
- Shared open files
 - stdout is the same in both parent and child

Modeling fork with Process Graphs

processes running concurrently

- A *process graph* is a useful tool for capturing the partial ordering of statements in a concurrent program:

- Each vertex is the execution of a statement
- $a \rightarrow b$ means a happens before b
- Edges can be labeled with current value of variables
- `printf` vertices can be labeled with output
- Each graph begins with a vertex with no incoming edges

- Any *topological sort* of the graph corresponds to a feasible total ordering.
 - Total ordering of vertices where all edges point from left to right

Process Graph Example

```


int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}


```

fork.c




Interpreting Process Graphs

Original graph:




Left-to-right ✓ Correct

Relabeled graph:




Feasible total ordering:



(viable)


Infeasible total ordering:



fork Example: Two consecutive forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

forks.c



Feasible output:

L0
L1
Bye
Bye
L1
Bye
Bye
Bye


Infeasible output:

L0
Bye ~ Bye happens after L1
L1
Bye
L1
Bye
Bye

fork Example: Nested forks in parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

forks.c



Feasible output:

L0
L1
Bye
Bye
L2
Bye

Infeasible output:

L0
Bye
L1
Bye
Bye
L2

La ghaff se fdece


fork Example: Nested forks in children

```

void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}

```

forks.c




Feasible output:

L0
Bye } order can
L1 } be different
L2 > Bye
Bye

Infeasible output:

L0
Bye
L1
Bye) L2 should be
Bye present
L2

Reaping Child Processes



■ Idea child

- When process terminates, it still consumes system resources
 - Examples: Exit status, various OS tables
 - Called a “zombie”
 - Living corpse, half alive and half dead

use these

some info in OS

- Reaping \rightsquigarrow only by parent

- Performed by parent on terminated child (using `wait` or `waitpid`)
 - Parent is given exit status information `exit(0)`
 - Kernel then deletes zombie child process

Parent waits for a process to die

■ What if parent doesn't reap?

- If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (pid == 1) *running in the system at the time*
 - So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

```
linux> ./forks 7 &
[1] 6639
```

```
void fork7() {
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n", getpid());
        exit(0);
    } else { parent
        printf("Running Parent, PID = %d\n", getpid());
        while (1)
            ; /* Infinite loop */ } }
```

parent infinite loop
never terminate
forks.c

```
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6639	ttyp9	00:00:03	forks
6640	ttyp9	00:00:00	forks <defunct>
6641	ttyp9	00:00:00	ps

```
linux> kill 6639
```

manually kill parent process

```
[1] Terminated
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6642	ttyp9	00:00:00	ps

parent process

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

■ **ps** shows child process as “defunct” (i.e., a zombie)

■ Killing parent allows child to be reaped by **init**

not reaping child process (done by init)

Non-terminating Child Example

nobody kills running child process

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6676 tttyp9      00:00:06 forks
 6677 tttyp9      00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6678 tttyp9      00:00:00 ps
```

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
               getpid());
        while (1)
            ; /* Infinite loop */
    } else { // Parent is gone
        printf("Terminating Parent, PID = %d\n",
               getpid());
        exit(0);
    }
}
```

forks.c

- Child process still active even though parent has terminated

- Must kill child explicitly, or else will keep running indefinitely

Such inaccuracies can kill your system
 parent process → zombie process

wait: Synchronizing with Children

explicitly

- Parent reaps a child by calling the `wait` function

- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the pid of the child process that terminated
 - If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`,
`WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`,
`WIFCONTINUED`
 - See textbook for details



wait: Synchronizing with Children

aligning steps
matching behavior

```

void fork9() {
    int child_status;
    if (fork() == 0) { child
        printf("HC: hello from child\n");
        exit(0);
    } else { parent
        printf("HP: hello from parent\n");
        wait(&child_status); sleep
        then if child_terminated() printf("CT: child has terminated\n");
        then wake up
        printf("Bye\n");
    }
}

```


forks.c

8 help

(P8)

→ P8 - &P &

she'll wait until
P8 finishes its job



Waiting for child process
to terminate

Feasible output:

- HC
- HP
- CT
- Bye

Infeasible output:

- HP
- CT
- Bye
- HC

Another wait Example

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```

void fork10() {
    pid_t pid[N];
    int i, child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0) {
            exit(100+i); /* Child */
        }
    for (i = 0; i < N; i++) { /* Parent */
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
normal exit status ↗ wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}

```

forks.c

PTD of child

for only 1 child to be completed

Parent does not get into loop body

obtain exit status from child status

actual status

waitpid: Waiting for a Specific Process terminated

pass process ID to wait for specific child process to be terminated

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - Suspends current process until specific process terminates
 - Various options (see textbook)

```

void fork11() {
    pid_t pid[N];
    int i;
    int child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = N-1; i >= 0; i--) {
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                   wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}

```

forkwards

Partly created child processes to be terminated

forks.c

System call execve: Loading and Running Programs

Want to create process of our own executable?

Fork → creates copy (if/else)

- int execve(char *filename, char *argv[], char *envp[])
- Loads and runs in the current process:

environment variable

- Executable file **filename**
 - Can be object file or script file beginning with #! interpreter (e.g., #!/bin/bash) → treat as executable
- ...with argument list **argv**
 - By convention **argv[0]==filename**
- ...and environment variable list **envp**
 - “name=value” strings (e.g., USER=droh)
 - getenv, putenv, printenv


Overwrites code, data, and stack *with your exec.*

- Retains PID, open files and signal context

running your exec from shell

Called once and never returns

- ...except if there is an error




Structure of the stack when a new program starts

VML of new program
space

execve will
overwrite
arguments

argv
(in %rsi)


argc
(in %rdi)



execve Example

included as default

- Executes “/bin/ls -lt /usr/include” in child process using current environment:



why both ~> maybe after fork, I would initialize both before execve

```

if ((pid = Fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}

```

Summary

standard control flow → jump, loop, if/else
inside application

Exceptions

- Events that require nonstandard control flow *(System event)*
- Generated externally (interrupts) or internally (traps and faults)

traps

Processes

- At any given time, system has multiple active processes
- Only one can execute at a time on a single core, though
- Each process appears to have total control of processor + private memory space

) illusion of each process

Summary (cont.)

- **Spawning processes**
 - Call fork *→ duplicate except some memory space*
 - One call, two returns
- **Process completion**
 - Call exit
 - One call, no return
- **Reaping and waiting for processes**
 - Call wait or waitpid *Cleanup messages*
- **Loading and running programs**
 - Call execve (or variant)
 - One call, (normally) no return