

Concurrent Programming

CS230: System Programming
17th Lecture

Instructor:
Jongse Park

Concurrent Programming is Hard!


- The human mind tends to be sequential
- The notion of time is often misleading
- Thinking about all possible sequences of events in a computer system is at least error prone and frequently impossible

Concurrent Programming is Hard!

- Classical problem classes of concurrent programs:
 - ***Races***: outcome depends on arbitrary scheduling decisions elsewhere in the system
 - Example: who gets the last seat on the airplane?
 - ***Deadlock***: improper resource allocation prevents forward progress
 - Example: traffic gridlock
 - ***Livelock / Starvation / Fairness***: external events and/or system scheduling decisions can prevent sub-task progress
 - Example: people always jump in front of you in line
- Many aspects of concurrent programming are beyond the scope of 15-213


Iterative Servers

- Iterative servers process one request at a time




Where Does Second Client Block?

- Second client attempts to connect to iterative server



- Call to `connect` returns
 - Even though connection not yet accepted
 - Server side TCP manager queues request
 - Feature known as “TCP listen backlog”
- Call to `rio_written` returns
 - Server side TCP manager buffers input data
- Call to `rio_readlineb` blocks
 - Server hasn't written anything for it to read yet.

Fundamental Flaw of Iterative Servers




- Solution: use *concurrent servers* instead
 - Concurrent servers use multiple concurrent flows to serve multiple clients at the same time

Creating Concurrent Flows

- Allow server to handle multiple clients simultaneously
- 1. Processes
 - Kernel automatically interleaves multiple logical flows
 - Each flow has its own private address space
- 2. Threads
 - Kernel automatically interleaves multiple logical flows
 - Each flow shares the same address space
- 3. I/O multiplexing with `select()`
 - Programmer manually interleaves multiple logical flows
 - All flows share the same address space
 - Relies on lower-level system abstractions

Concurrent Servers: Multiple Processes

- Spawn separate process for each client



Review: Iterative Echo Server

```
int main(int argc, char **argv)
{
    int listenfd, connfd;
    int port = atoi(argv[1]);
    struct sockaddr_in clientaddr;
    int clientlen = sizeof(clientaddr);

    listenfd = open_listenfd(port);
    while (1) {
        connfd = accept(listenfd, (SA *) &clientaddr, &clientlen);
        echo(connfd);
        close(connfd);
    }
    exit(0);
}
```

- Accept a connection request
- Handle echo requests until client terminates

Process-Based Concurrent Server

```
int main(int argc, char **argv)
{
    int listenfd, connfd;
    int port = atoi(argv[1]);
    struct sockaddr_in clientaddr;
    int clientlen=sizeof(clientaddr);

    signal(SIGCHLD, sigchld_handler);
    listenfd = open_listenfd(port);
    while (1) {
        connfd = accept(listenfd, (SA *) &clientaddr, &clientlen);
        if (fork() == 0) {
            close(listenfd); /* Child closes its listening socket */
            echo(connfd);    /* Child services client */
            close(connfd);   /* Child closes connection with client */
            exit(0);          /* Child exits */
        }
        Close(connfd); /* Parent closes connected socket (important!) */
    }
}
```


Fork separate process for each client
Does not allow any communication between different client handlers

Process-Based Concurrent Server (cont)

```
void sigchld_handler(int sig)
{
    while (waitpid(-1, 0, WNOHANG) > 0)
    ;
    return;
}
```

- Reap all zombie children

Process Execution Model




- Each client handled by independent process
- No shared state between them
- Both parent & child have copies of listenfd and connfd
 - Parent must close connfd
 - Child must close listenfd


Concurrent Server: accept Illustrated



- 1. Server blocks in accept, waiting for connection request on listening descriptor listenfd*



- 2. Client makes connection request by calling and blocking in connect*



- 3. Server returns connfd from accept. Forks child to handle client. Client returns from connect. Connection is now established between clientfd and connfd*

Implementation Must-dos With Process-Based Designs

- Listening server process must reap zombie children
 - to avoid fatal memory leak
- Listening server process must close its copy of `connfd`
 - Kernel keeps reference for each socket/open file
 - After `fork`, `refcnt(connfd) = 2`
 - Connection will not be closed until `refcnt(connfd) == 0`

View from Server's TCP Manager

Client 1 Client 2 Server

```
srv> ./echoserverp 15213
```

```
cl1> ./echoclient greatwhite.ics.cs.cmu.edu 15213
```

```
srv> connected to (128.2.192.34), port 50437
```

```
cl2> ./echoclient greatwhite.ics.cs.cmu.edu 15213
```

```
srv> connected to (128.2.205.225), port 41656
```

Connection	Host	Port	Host	Port
Listening	---	---	128.2.220.10	15213
cl1	128.2.192.34	50437	128.2.220.10	15213
cl2	128.2.205.225	41656	128.2.220.10	15213

View from Server's TCP Manager

Connection	Host	Port	Host	Port
Listening	---	---	128.2.220.10	15213
c11	128.2.192.34	50437	128.2.220.10	15213
c12	128.2.205.225	41656	128.2.220.10	15213

- Port Demultiplexing
 - TCP manager maintains separate stream for each connection
 - Each represented to application program as socket
 - New connections directed to listening socket
 - Data from clients directed to one of the connection sockets

Pros and Cons of Process-Based Designs


- + Handle multiple connections concurrently
- + Clean sharing model
 - descriptors (no)
 - file tables (yes)
 - global variables (no)
- + Simple and straightforward
- – Additional overhead for process control
- – Nontrivial to share data between processes
 - Requires IPC (interprocess communication) mechanisms
 - FIFO's (named pipes), shared memory, and semaphores

Approach #2: Multiple Threads

- Very similar to approach #1 (multiple processes)
 - but, with threads instead of processes


Traditional View of a Process

- Process = process context + code, data, and stack



Alternate View of a Process

- Process = thread + code, data, and kernel context



A Process With Multiple Threads


- Multiple threads can be associated with a process
 - Each thread has its own logical control flow
 - Each thread shares the same code, data, and kernel context
 - Share common virtual address space (inc. stacks)
 - Each thread has its own thread id (TID)

Thread 1 (main thread)



Thread 1 context:
Data registers
Condition codes
SP1
PC1

Shared code and data



Kernel context:
VM structures
Descriptor table
brk pointer

Thread 2 (peer thread)




Thread 2 context:
Data registers
Condition codes
SP2
PC2


Logical View of Threads

- Threads associated with process form a pool of peers
 - Unlike processes which form a tree hierarchy

Threads associated with process foo





Process hierarchy




Thread Execution

- Single Core Processor
 - Simulate concurrency by time slicing
- Multi-Core Processor
 - Can have true concurrency



Logical Concurrency

- Two threads are (logically) concurrent if their flows overlap in time
- Otherwise, they are sequential
- Examples:
 - Concurrent: A & B, A&C
 - Sequential: B & C



Threads vs. Processes

- How threads and processes are similar
 - Each has its own logical control flow
 - Each can run concurrently with others (possibly on different cores)
 - Each is context switched
- How threads and processes are different
 - Threads share code and some data
 - Processes (typically) do not
 - Threads are somewhat less expensive than processes
 - Process control (creating and reaping) is twice as expensive as thread control
 - Linux numbers:
 - ~20K cycles to create and reap a process
 - ~10K cycles (or less) to create and reap a thread

Posix Threads (Pthreads) Interface

- *Pthreads*: Standard interface for ~60 functions that manipulate threads from C programs
 - Creating and reaping threads
 - `pthread_create()`
 - `pthread_join()`
 - Determining your thread ID
 - `pthread_self()`
 - Terminating threads
 - `pthread_cancel()`
 - `pthread_exit()`
 - `exit()` [terminates all threads] , `RET` [terminates current thread]
 - Synchronizing access to shared variables
 - `pthread_mutex_init`
 - `pthread_mutex_[un]lock`
 - `pthread_cond_init`
 - `pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```
/*
 * hello.c - Pthreads "hello, world" program
 */
#include "csapp.h"

void *thread(void *vargp);

int main() {
    pthread_t tid;
    Pthread_create(&tid, NULL, thread, NULL);
    Pthread_join(tid, NULL);
    exit(0);
}
```


```
/* thread routine */
void *thread(void *vargp) {
    printf("Hello, world!\n");
    return NULL;
}
```

*Thread attributes
(usually NULL)*

*Thread arguments
(void *p)*

*return value
(void **p)*

Execution of Threaded “hello, world”



Thread-Based Concurrent Echo Server

```
int main(int argc, char **argv) {
    int port = atoi(argv[1]);
    struct sockaddr_in clientaddr;
    int clientlen = sizeof(clientaddr);
    pthread_t tid;

    int listenfd = open_listenfd(port);
    while (1) {
        int *connfdp = malloc(sizeof(int));
        *connfdp = accept(listenfd,
                           (SA *) &clientaddr, &clientlen);
        pthread_create(&tid, NULL, echo_thread, connfdp);
    }
}
```


- Spawn new thread for each client
- Pass it copy of connection file descriptor
- Note use of Malloc()
 - Without corresponding Free()

Thread-Based Concurrent Server (cont)

```
/* thread routine */
void *echo_thread(void *vargp)
{
    int connfd = *((int *)vargp);
    pthread_detach(pthread_self());
    free(vargp);
    echo(connfd);
    close(connfd);
    return NULL;
}
```

- Run thread in “detached” mode
 - Runs independently of other threads
 - Reaped when it terminates
- Free storage allocated to hold clientfd
 - “Producer-Consumer” model

Threaded Server Execution Model



- Each client handled by individual peer thread
- Threads share all process state except TID
- Each thread has a separate stack for local variables

Issues With Thread-Based Servers

- Must run “detached” to avoid memory leak
 - At any point in time, a thread is either *joinable* or *detached*
 - *Joinable* thread can be reaped and killed by other threads
 - must be reaped (with `pthread_join`) to free memory resources
 - *Detached* thread cannot be reaped or killed by other threads
 - resources are automatically reaped on termination
 - Default state is joinable
 - use `pthread_detach(pthread_self())` to make detached
- Must be careful to avoid unintended sharing
 - For example, passing pointer to main thread’s stack
 - `Pthread_create(&tid, NULL, thread, (void *)&connfd);`
- All functions called by a thread must be *thread-safe*
 - (next lecture)

Potential Form of Unintended Sharing

```

while (1) {
    int connfd = accept(listenfd, (SA *) &clientaddr, &clientlen);
    pthread_create(&tid, NULL, echo_thread, (void *) &connfd);
}
}

```

main thread

Main thread stack

`connfd = connfd1`

`connfd`

`connfd = connfd2`

`connfd = *vargp`

`peer1`

Race!

`peer2`

Peer₁ stack

`vargp`

Peer₂ stack

`vargp`

Why would both copies of vargp point to same location?

Could this race occur?

Main

```
int i;
for (i = 0; i < 100; i++) {
    pthread_create(&tid, NULL,
                  thread, &i);
}
```

Thread


```
void *thread(void *vargp)
{
    int i = *((int *)vargp);
    pthread_detach(pthread_self());
    save_value(i);
    return NULL;
}
```

■ Race Test


- If no race, then each thread would get different value of i
- Set of saved values would consist of one copy each of 0 through 99.

Experimental Results


No Race



Single core laptop



Multicore server



- The race can really happen!

Pros and Cons of Thread-Based Designs

- + Easy to share data structures between threads
 - e.g., logging information, file cache.
- + Threads are more efficient than processes.

- – Unintentional sharing can introduce subtle and hard-to-reproduce errors!
 - The ease with which data can be shared is both the greatest strength and the greatest weakness of threads.
 - Hard to know which data shared & which private
 - Hard to detect by testing
 - Probability of bad race outcome very low
 - But nonzero!
 - Future lectures

Approaches to Concurrency

- Processes
 - Hard to share resources: Easy to avoid unintended sharing
 - High overhead in adding/removing clients
- Threads
 - Easy to share resources: Perhaps too easy
 - Medium overhead
 - Not much control over scheduling policies
 - Difficult to debug
 - Event orderings not repeatable
- I/O Multiplexing
 - Tedious and low level
 - Total control over scheduling
 - Very low overhead
 - Cannot create as fine grained a level of concurrency
 - Does not make use of multi-core