

Synchronization: Basics

CS230: System Programming
18th Lecture

Instructor:
Jongse Park

Today

- Sharing
- Mutual exclusion
- Semaphores

Shared Variables in Threaded C Programs

- Question: Which variables in a threaded C program are shared?
 - The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”
- Requires answers to the following questions:
 - What is the memory model for threads?
 - How are instances of variables mapped to memory?
 - How many threads might reference each of these instances?
- *Def:* A variable x is *shared* if and only if multiple threads reference some instance of x .

Threads Memory Model


- Conceptual model:
 - Multiple threads run within the context of a single process
 - Each thread has its own separate thread context
 - Thread ID, stack, stack pointer, PC, condition codes, and GP registers
 - All threads share the remaining process context
 - Code, data, heap, and shared library segments of the process virtual address space
 - Open files and installed handlers
- Operationally, this model is not strictly enforced:
 - Register values are truly separate and protected, but...
 - Any thread can read and write the stack of any other thread

*The mismatch between the conceptual and operation model
is a source of confusion and errors*

Example Program to Illustrate Sharing

```
char **ptr; /* global */  
  
int main()  
{  
    int i;  
    pthread_t tid;  
    char *msgs[2] = {  
        "Hello from foo",  
        "Hello from bar"  
    };  
    ptr = msgs;  
  
    for (i = 0; i < 2; i++)  
        Pthread_create(&tid,  
                       NULL,  
                       thread,  
                       (void *)i);  
    pthread_exit(NULL);  
}
```

```
/* thread routine */  
void *thread(void *vargp)  
{  
    int myid = (int) vargp;  
    static int cnt = 0;  
  
    printf("[%d]: %s (svar=%d)\n",  
           myid, ptr[myid], ++cnt);  
}
```



*Peer threads reference main thread's stack
indirectly through global ptr variable*

Mapping Variable Instances to Memory

- Global variables
 - *Def*: Variable declared outside of a function
 - **Virtual memory contains exactly one instance of any global variable**
- Local variables
 - *Def*: Variable declared inside function without `static` attribute
 - **Each thread stack contains one instance of each local variable**
- Local static variables
 - *Def*: Variable declared inside function with the `static` attribute
 - **Virtual memory contains exactly one instance of any local static variable.**

Mapping Variable Instances to Memory

Global var: 1 instance (ptr [data])

```
char **ptr; /* global */  
  
int main()  
{  
    int i;  
    pthread_t tid;  
    char *msgs[2] = {  
        "Hello from foo",  
        "Hello from bar"  
    };  
    ptr = msgs;  
  
    for (i = 0; i < 2; i++)  
        pthread_create(&tid,  
                      NULL,  
                      thread,  
                      (void *)i);  
    pthread_exit(NULL);  
}
```

Local vars: 1 instance (i.m, msgs.m)

Local var: 2 instances (
myid.p0 [peer thread 0's stack],
myid.p1 [peer thread 1's stack]
)

```
/* thread routine */  
void *thread(void *vargp)  
{  
    int myid = (int)vargp;  
    static int cnt = 0;  
  
    printf("[%d]: %s (svar=%d)\n",  
          myid, ptr[myid], ++cnt);  
}
```

Local static var: 1 instance (cnt [data])

Shared Variable Analysis

- Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?peer thread 1?</i>	
ptr	yes	yes	yes
cnt	no	yes	yes
i.m	yes	no	no
msgs.m	yes	yes	yes
myid.p0	no	yes	no
myid.p1	no	no	yes

- Answer: A variable x is shared iff multiple threads reference at least one instance of x . Thus:
 - **ptr, cnt, and msgs are shared**
 - **i and myid are *not* shared**

Today

- Sharing
- Mutual exclusion
- Semaphores

badcnt.c: Improper Synchronization

```

volatile int cnt = 0; /* global */

int main(int argc, char **argv)
{
    int niters = atoi(argv[1]);
    pthread_t tid1, tid2;

    pthread_create(&tid1, NULL,
                  thread, &niters);
    pthread_create(&tid2, NULL,
                  thread, &niters);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%d\n", cnt);
    else
        printf("OK cnt=%d\n", cnt);
    exit(0);
}

```

```

/* Thread routine */
void *thread(void *vargp)
{
    int i, niters = *((int *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}

```

```

linux> ./badcnt 10000
OK cnt=20000
linux> ./badcnt 10000
BOOM! cnt=13051
linux>

```

cnt should equal 20,000.

What went wrong?

Assembly Code for Counter Loop

C code for counter loop in thread i

```
for (i=0; i < niters; i++)
    cnt++;
```

Corresponding assembly code

```
        movl (%rdi), %ecx
        movl $0, %edx
        cmpl %ecx, %edx
        jge .L13
.L11:
        movl cnt(%rip), %eax
        incl %eax
        movl %eax, cnt(%rip)
        incl %edx
        cmpl %ecx, %edx
        jl .L11
.L13:
```

The assembly code is annotated with curly braces on the right side, grouping it into four sections:


- Head (H_i)**: The first four instructions: `movl (%rdi), %ecx`, `movl $0, %edx`, `cmpl %ecx, %edx`, and `jge .L13`.
- Load cnt (L_i)**: The instruction `movl cnt(%rip), %eax`.
- Update cnt (U_i)**: The instruction `incl %eax`.
- Store cnt (S_i)**: The instruction `movl %eax, cnt(%rip)`.

Below the `Store` section, there is another brace grouping the `incl %edx`, `cmpl %ecx, %edx`, and `jl .L11` instructions under the label **Tail (T_i)**.

Concurrent Execution

- **Key idea:** In general, any sequentially consistent interleaving is possible, but some give an unexpected result!
 - I_i denotes that thread i executes instruction I
 - $\%eax_i$ is the content of $\%eax$ in thread i 's context

i (thread)	$instr_i$	$\%eax_1$	$\%eax_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2

 Thread 1
critical section

 Thread 2
critical section

OK

Concurrent Execution (cont)

- Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2

i (thread)	instr _i	%eax ₁	%eax ₂	cnt
1	H ₁	-	-	0
1	L ₁	0	-	0
1	U ₁	1	-	0
2	H ₂	-	-	0
2	L ₂	-	0	0
1	S ₁	1	-	1
1	T ₁	1	-	1
2	U ₂	-	1	1
2	S ₂	-	1	1
2	T ₂	-	1	1

Oops!

Concurrent Execution (cont)

- How about this ordering?


i (thread)	instr _i	%eax ₁	%eax ₂	cnt
1	H ₁			0
1	L ₁	0		
2	H ₂			
2	L ₂		0	
2	U ₂		1	
2	S ₂		1	1
1	U ₁	1		
1	S ₁	1		1
1	T ₁			
2	T ₂			1

Oops!

- We can analyze the behavior using a *progress graph*

Progress Graphs

Thread 2



A *progress graph* depicts the discrete *execution state space* of concurrent threads.


Each axis corresponds to the sequential order of instructions in a thread.

Each point corresponds to a possible *execution state* $(Inst_1, Inst_2)$.

E.g., (L_1, S_2) denotes state where thread 1 has completed L_1 and thread 2 has completed S_2 .

Trajectories in Progress Graphs

Thread 2




A *trajectory* is a sequence of legal state transitions that describes one possible concurrent execution of the threads.

Example:

$H_1, L_1, U_1, H_2, L_2, S_1, T_1, U_2, S_2, T_2$

Critical Sections and Unsafe Regions




L , U , and S form a *critical section* with respect to the shared variable `cnt`

Instructions in critical sections (wrt to some shared variable) should not be interleaved

Sets of states where such interleaving occurs form *unsafe regions*

Critical Sections and Unsafe Regions

Thread 2



Def: A trajectory is **safe** iff it does not enter any unsafe region

Claim: A trajectory is correct (wrt cnt) iff it is safe

Enforcing Mutual Exclusion

- *Question:* How can we guarantee a safe trajectory?
- Answer: We must **synchronize** the execution of the threads so that they never have an unsafe trajectory.
 - i.e., need to guarantee **mutually exclusive access** to critical regions
- Classic solution:
 - Semaphores (Edsger Dijkstra)
- Other approaches (out of our scope)
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

Today

- Sharing
- Mutual exclusion
- Semaphores

Semaphores

- **Semaphore:** non-negative global integer synchronization variable
- Manipulated by *P* and *V* operations:
 - *P(s)*: [**while** (*s* == 0) **wait()**; *s*--;]
 - Dutch for "Proberen" (test)
 - *V(s)*: [*s*++;]
 - Dutch for "Verhogen" (increment)
- OS kernel guarantees that operations between brackets [] are executed indivisibly
 - Only one *P* or *V* operation at a time can modify *s*.
 - When **while** loop in *P* terminates, only that *P* can decrement **s**
- **Semaphore invariant: (*s* >= 0)**

C Semaphore Operations

Pthreads functions:

```
#include <semaphore.h>

int sem_init(sem_t *sem, 0, unsigned int val); /* s = val */
int sem_wait(sem_t *s); /* P(s) */
int sem_post(sem_t *s); /* V(s) */
```

CS:APP wrapper functions:

```
#include "csapp.h"

void P(sem_t *s); /* Wrapper function for sem_wait */
void V(sem_t *s); /* Wrapper function for sem_post */
```

badcnt.c: Improper Synchronization

```
volatile int cnt = 0; /* global */

int main(int argc, char **argv)
{
    int niters = atoi(argv[1]);
    pthread_t tid1, tid2;

    pthread_create(&tid1, NULL,
                  thread, &niters);
    pthread_create(&tid2, NULL,
                  thread, &niters);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%d\n", cnt);
    else
        printf("OK cnt=%d\n", cnt);
    exit(0);
}
```

```
/* Thread routine */
void *thread(void *vargp)
{
    int i, niters = *((int *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}
```

How can we fix this using semaphores?

Using Semaphores for Mutual Exclusion

- Basic idea:
 - Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables).
 - Surround corresponding critical sections with $P(\text{mutex})$ and $V(\text{mutex})$ operations.
- Terminology:
 - *Binary semaphore*: semaphore whose value is always 0 or 1
 - *Mutex*: binary semaphore used for mutual exclusion
 - P operation: “*locking*” the mutex
 - V operation: “*unlocking*” or “*releasing*” the mutex
 - “*Holding*” a mutex: locked and not yet unlocked.
 - *Counting semaphore*: used as a counter for set of available resources.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable cnt :

```
volatile int cnt = 0;      /* Counter */  
sem_t mutex;              /* Semaphore that protects cnt */  
  
sem_init(&mutex, 0, 1);   /* mutex = 1 */
```

- Surround critical section with P and V:


```
for (i = 0; i < niters; i++) {  
    P(&mutex);  
    cnt++;  
    V(&mutex);  
}
```

```
linux> ./goodcnt 10000  
OK cnt=20000  
linux> ./goodcnt 10000  
OK cnt=20000  
linux>
```

Warning: It's much slower than
badcnt.c.

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

Semaphore invariant creates a *forbidden region* that encloses unsafe region that cannot be entered by any trajectory.

Summary

- Programmers need a clear model of how variables are shared by threads.
- Variables shared by multiple threads must be protected to ensure mutually exclusive access.
- Semaphores are a fundamental mechanism for enforcing mutual exclusion.