

CS330 Operating Systems and Lab.

Virtual Memory

Spring 2026



Instructor : Insik Shin

Virtual Memory

- Virtual memory
 - A technique to separate user logical memory from physical memory
 - It allows
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Address spaces can be shared by several processes
 - More efficient process creation
- Virtual memory can be implemented via:
 - Demand paging

Demand Paging

- Demand paging
 - A paging system that brings a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users

Valid-Invalid Bit

- Some pages can be loaded into memory, but the others not
 - Need some form of hardware support to distinguish which pages are in memory / on disk
- Valid-invalid bit in page table
 - With each page table entry, a valid–invalid bit is associated (**v** $\Rightarrow$ in-memory, **i** $\Rightarrow$ not-in-memory)
 - Initially valid–invalid bit is set to **i** on all entries
 - During address translation, if valid–invalid bit in page table entry is **i** $\Rightarrow$ page fault

Frame # valid-invalid bit

	v
	i
....	
	i

page table

Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical
memory

valid-invalid bit		
frame		
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table


0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory



Page Fault

- If a process tries to access a page that was not brought into memory,
 1. Paging hardware notices that the validation bit = **i** and causes a trap to OS
 2. OS looks at an internal table (in PCB) to decide:
 - Invalid reference $\Rightarrow$ abort
 - Just not in memory
 3. Get empty frame
 4. Swap page into frame
 5. Reset page tables (i.e., set validation bit = **v**)
 6. Restart the instruction that caused the page fault



Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$, no page faults
 - if $p = 1$, every reference is a fault
- Effective Access Time (EAT)
EAT = $(1 - p) \times \text{memory access}$
+ $p (\text{page fault overhead} + \text{memory access})$

page fault overhead = swap page out + swap page in
+ restart overhead


Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p (8 \text{ milliseconds} + 200) \\ &= (1 - p) \times 200 + p \times 8,000,200 \\ &= 200 + p \times 8,000,000 \end{aligned}$$
- If **one access out of 1,000** causes a page fault, then $\text{EAT} = 8.2 \text{ microseconds}$.
This is a slowdown by a factor of 40!!

Demand Paging

- Why does this work?
 - *Locality of reference*
 - **Temporal locality**: data elements are frequently referenced again within relatively small time durations
 - **Spatial locality**: data elements within relatively close storage locations are frequently referenced
- Locality implies infrequent paging !


Locality In A Memory-Reference Pattern



Page Fault (revisit)

- Page fault happens if a process tries to access a page that was not brought into memory
 1. Operating system looks at another table to decide:
 - Invalid reference $\Rightarrow$ abort
 - Just not in memory $\Rightarrow$ trap
 2. Get empty frame
 3. Swap page into frame
 4. Reset page tables (i.e., set validation bit = **v**)
 5. Restart the instruction that caused the page fault


Steps in Handling a Page Fault




What happens if there is no free frame?

- Page replacement – find some page in memory, but not really in use, swap it out
 - algorithm
 - performance – want an algorithm which will result in minimum number of page faults
- Same pages may be brought into memory several times

Need For Page Replacement



Page Replacement



Page Replacement Algorithms

- Select a victim page to swap out, aiming at a lower page-fault rate
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)


1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

- 4 frames


1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

- Belady's Anomaly: more frames $\Rightarrow$ more page faults

Graph of Page Faults vs. the Number of Frames



FIFO Illustrating Belady's Anomaly



Optimal Algorithm

- Replace page that **will not be used** for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1
2
3
4

4

6 page faults

5

- How do you know this?
- Used for measuring how well your algorithm performs

Least Recently Used (LRU) Algorithm

- Replace the page that *has not been used* for the longest time
- Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**


1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

LRU Algorithm (Cont.)

- Stack implementation – keep a stack of page numbers in a double link form:
 - Page referenced:
 - move it to the top; requires 6 pointers to be changed
 - No search for replacement

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack before “a”




stack before “b”




Use Of A Stack to Record The Most Recent Page References

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2



stack
before
a



stack
after
b



Stack Algorithms

- Interesting feature
 - Stack algorithms, such as optimal and LRU, can never exhibit Belady's anomaly
- Stack algorithm
 - A set of pages in memory for n frames is always a *subset* of the set of pages in memory for $n+1$ frames

- 4 frames



- 5 frames



Least Recently Used (LRU) Algorithm

- Replace the page that *has not been used* for the longest time
- Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

- Counter/clock implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters to determine which are to change


LRU-Approximation Algorithms

- Background
 - Little hardware support for true LRU, but some h/w help available for LRU approximation
- Additional-reference-bits algorithm
 - Take advantage of the reference bit of PTE
 - At regular intervals, record the reference bit for each page to figure out later which page has been used less recently


LRU-Approximation Algorithms

- Second-chance algorithm
 - Need reference bit
 - Clock replacement (or FIFO replacement)
 - If page to be replaced (in clock order) has reference bit = 1 then:
 - set reference bit 0
 - leave page in memory
 - replace next page (in clock order), subject to same rules

Second-Chance (clock) Page-Replacement Algorithm



(a)




(b)

Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- **LFU Algorithm**: replaces page with smallest count
- **MFU Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Steps in Handling a Page Fault



CPU is idle during page fault handling

Page Replacement (Memory Reclamation)

frame valid-invalid bit


0	i
f	v

page table

② change to invalid

④ reset page table for new page


How many processes (how big logical addresses) can be supported ?



Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - low CPU utilization
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system
- **Thrashing** $\equiv$ a process is spending much time in paging (maybe more time than executing)


Thrashing (Cont.)



Demand Paging and Thrashing

- Why does demand paging work?
 - Locality model
 - Process migrates from one locality to another
 - Localities may overlap
- Why does thrashing occur?
 - Σ size of locality > total memory size

Locality In A Memory-Reference Pattern




Working Set Model


- Working set window (Δ)
 - a fixed number of page references (i.e., 10,000 instruction)
- Working set
 - A set of pages that a process references in the last Δ
 - Approximation of program's locality
- Working set size
 - A total number of pages in the working set (varies in time)

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

Working Set Model


- Working set window (Δ)
 - if Δ is too small, it will not encompass entire locality
 - if Δ is too large, it will encompass several localities
 - if $\Delta = \infty$, then it will encompass entire program
- Working-set strategy can prevent thrashing
 - $D = \sum WSS_i \equiv$ total demand frames
 - if $D > m \Rightarrow$ Thrashing (m : total number of available frames)
 - Policy if $D > m$, then suspend one of the processes

Page-Fault Frequency Scheme


- Thrashing handling
 - Needs to detect thrashing
 - Working-set model works ($D > m$); but clumsy
 - Page-fault frequency scheme is more direct
- Page fault frequency scheme
 - Thrashing has a high page-fault rate

Page-Fault Frequency Scheme

- Establish “acceptable” page-fault rate
 - If actual rate too low, process loses frame
 - If actual rate too high, process gains frame



Paging Model of Logical and Physical Memory




Process Creation

- Virtual memory allows other benefits during process creation:
 - Memory Sharing
 - Copy-on-Write
 - Memory-Mapped Files

Page Sharing

- Private virtual address space protect applications from each other
 - However, this makes it difficult to share data
- Paging makes sharing easier
 - Multiple processes can have their own logical addresses mapped to the same physical address
 - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).

Shared Pages Example




Copy-on-Write

- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory

If either process modifies a shared page, only then is the page copied


- COW allows more efficient process creation as only modified pages are copied
- Free pages are allocated from a **pool** of zeroed-out pages

Before Process 1 Modifies Page C



Suppose Process 1 wants to write something on page C


After Process 1 Modifies Page C



Memory-Mapped Files

- Consider a user program that reads a long file
 - as examples, editing documents, watching videos, etc.
- Two ways to read/write files
 - through **read()/write()** systems calls
 - through **memory-mapped files** (memory operations)
- Performance issues
 - system call vs. memory operations



Memory-Mapped Files




Memory-Mapped Files

- **Memory-mapped file I/O offers many benefits**
- **Performance benefits**
 - It allows file I/O to be treated as routine memory access (e.g., like an array) by **mapping** a disk block to a page in memory, offering significant performance benefits:
 - A file is initially read using demand paging.
 - Subsequent reads/writes to/from the file are performed through ordinary memory accesses (i.e., using pointers).
 - Using ordinary memory accesses is much cheaper than using **read()/write()** system calls
 - Closing the file results in all the memory-mapped data being written to disk and removed from the virtual memory
- **Easy memory sharing**
 - It also allows several processes to map the same file allowing the pages in memory to be shared

Memory Mapped Files



Paging Model of Logical and Physical Memory



Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
 - A segment is a logical unit such as:

main program

procedure

function

method

object


local variables, global variables

common block


stack

symbol table


arrays



Logical View of Segmentation



user space




physical memory space


Segmentation Architecture

- Logical address consists of a two tuple:
 <**segment-number, offset**> ,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program; segment number **s** is legal if **s < STLR**

Segmentation Hardware




Example of Segmentation




	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table




Example: The Intel Pentium


- Supports both segmentation and segmentation with paging
- CPU generates logical address
 - Given to segmentation unit
 - which produces linear addresses
 - Linear address given to paging unit
 - which generates physical address in main memory
 - paging units form equivalent of MMU




Logical address




Linear address



Physical address



Intel Pentium Segmentation



Pentium Paging Architecture

